

CYBER502x

Computer Forensics

Unit 3: Linux/Unix Filesystems

Investigating Linux/Unix systems

- Four basic forensics steps
 - Collect
 - Preserve
 - *Analyze*
 - Present (report)

What is special about forensics analysis?

- Show deleted content and other data that is typically inaccessible
- Parse data directly from the image bypassing kernel's support
- Work on both images and live systems

Lets talk about the filesystem...

- Recall the boot process...
- BIOS/MBR
 - BIOS starts POST
 - Finds DRIVES using BIOS sequence
 - Finds BOOTABLE MEDIA
 - MBR, located in sector zero, contains the initial boot code
- EFI (Extensible Firmware Interface)/GPT
 - EFI jumps to the EFI system partition (ESP) that begin the initial boot strapping

GUID Partition Table (GPT)

- GPT scheme and EFI boot framework
- Supports up to 128 primary partitions on a GPT disk
- Provides unique identification of both disks and partitions
- Tools for managing GPT
 - OSX – Disk Utility (or gpt)
 - Windows – DISKPART
 - Linux – parted /dev/xxx print
 - mmls -t gpt /dev/xxx

GPT Forensic Analysis (not required)

- Bruce J. Nikkel, *Forensic Analysis of GPT Disks and GUID Partition Tables*, The International Journal of Digital Forensics and Incident Response Vol. 6, No. 1-2
- Brian Carrier, *File System Forensic Analysis*, Addison Wesley, 2005

MBR structure

- MBR executable code starts at offset 0x0000, total 446 bytes
 - The MBR messages start at offset 0x008b
- The partition table starts at offset 0x01be
 - 64 bytes in four 16 bytes sections
 - 1st partition Entry starts at offset 0x01be
 - 2nd partition entry starts at offset 0x01ce
 - 3rd partition entry starts at offset 0x01de
 - 4th partition entry starts at offset 0x01ee
- The signature is at offset 0x01fe, 2 bytes (55AAh)
- Total 512 bytes

File systems in Unix/Linux

- Unix File System
 - UFS, or Solaris and BSD systems BSD derived (from Berkeley fast file system (FFS))
- Linux File System – derived from UFS
 - Extended file system (ext) in 1992
 - ext, ext2 (1993), ext3(2001), ext4 (2008)

Disk Structure

- Disks are partitioned, and filesystem is installed
- Each partition contains the following areas:
 - An optional boot block
 - A super block defining the boundaries of the other areas.
 - Block Bitmap
 - Inode Bitmap
 - A set of file information blocks known as I-nodes
 - The data blocks (free and used intermingled)

superblock

- Superblock contains
 - Magic Number
 - For EXT2, it is 0xEF53
 - Mount Count and Maximum Mount Count
 - Block Size, for example 4096B
 - Inode count and Block count
 - Number of free disk blocks
 - Number of free inodes on the system
 - First inode
 - This is the inode number of the first inode in the file system. The first inode in an EXT2 root file system would be the directory entry for the '/' directory


Inodes

- Contain META data (info about files)
 - type
 - Owner
 - Permission bits
 - MAC times
 - Links to the file (link count)
 - Data block addresses

Inode structure (each row is 32 bits)

Type/Mode	Link Count
User ID	Group ID
File Size (Bytes)	
Creation/Modification/Access Time	
Direct Block Pointer 1	
Direct Block Pointer 2	
.	
.	
.	
.	
.	
.	
.	
.	
.	
Direct Block Pointer 12	
Indirect Block Pointer	
Double Indirect Block Pointer	
Triple Indirect Block Pointer	

12 Direct Block Pointers



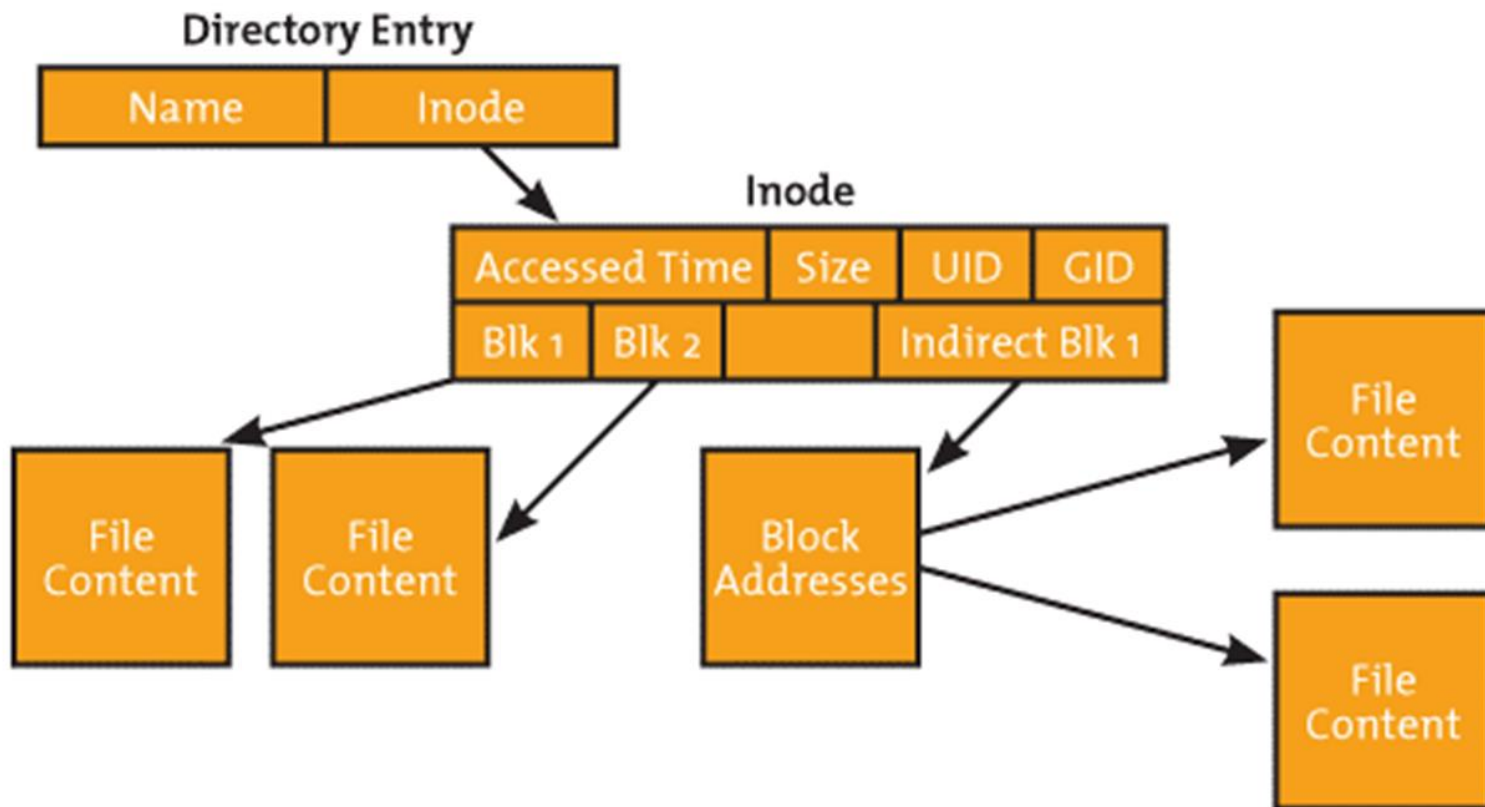
Directories

- Everything in Linux is a file
- Directory is a simple file whose data is a sequence of file entries which contains:
 - Inode number
 - Byte offset in directory (or the length of this entry)
 - File name

Directory example

Byte offset in directory	Inode number	File name
0	70	.
16	35	..
32	123	file1
48	345	file2
64	90	file3

Directory entry



What happened when we create a file

- Each file has data stored in **blocks, inodes and directory entries**
 - A free inode is chosen from inode bitmap
 - The superblock free-inode values are decremented
 - Add an entry in the parent directory
 - Choose a free data block from data bitmap to contain the file contents
- Fill in the inode contents

How files are deleted...

- <http://www.porcupine.org/forensics/forensic-discovery/chapter4.html>
- When the link-count in the inode reaches zero (0):
 - The Data blocks in the Block Bitmap are marked as free
 - The inode in the Inode Bitmap is marked as free
 - The deletion time is set in the inode.
 - The directory entry is marked as unused.

What happens when a file is deleted in ext3/ext4?

- The file size in the inode is set to zero.
- The data blocks info in the inode is cleared.
- <http://linux.sys-con.com/node/117909> (not required)
- *An analysis of Ext4 for digital forensics*, Kevin Fairbanks, 2012, in dfrws.org proceedings (not required)

Why do I care about deletion?

- Data still exists on disk
 - Recoverable until space is overwritten
 - Larger disks less likely to overwrite formerly-used space
- Therefore, you can (usually) recover deleted files
 - Unless the disk blocks are “wiped” before the file is deleted (e.g. “srm” or “shred”) (`dd if=/dev/zero` or `/dev/random`)